

Distributed ES with Akka Persistence

Duncan K. DeVore

SWE @ Typesafe

@ironfish



Outline

1. Introduction
2. What is Event Sourcing?
3. What are Commands?
4. What are Events?
5. What is Akka Persistence?
6. What is Akka Cluster Sharding?
7. Consistency
8. Conclusion

Outline

1. Introduction
2. What is Event Sourcing?
3. What are Commands?
4. What are Events?
5. What is Akka Persistence?
6. What is Akka Cluster Sharding?
7. Consistency
8. Conclusion

It's a different world
out there

Yesterday

Single machines

Single core processors

Expensive RAM

Expensive disk

Slow networks

Few concurrent users

Small data sets

Latency in seconds

Today

Clusters of machines

Multicore processors

Cheap RAM

Cheap disk

Fast networks

Lots of concurrent users

Large data sets

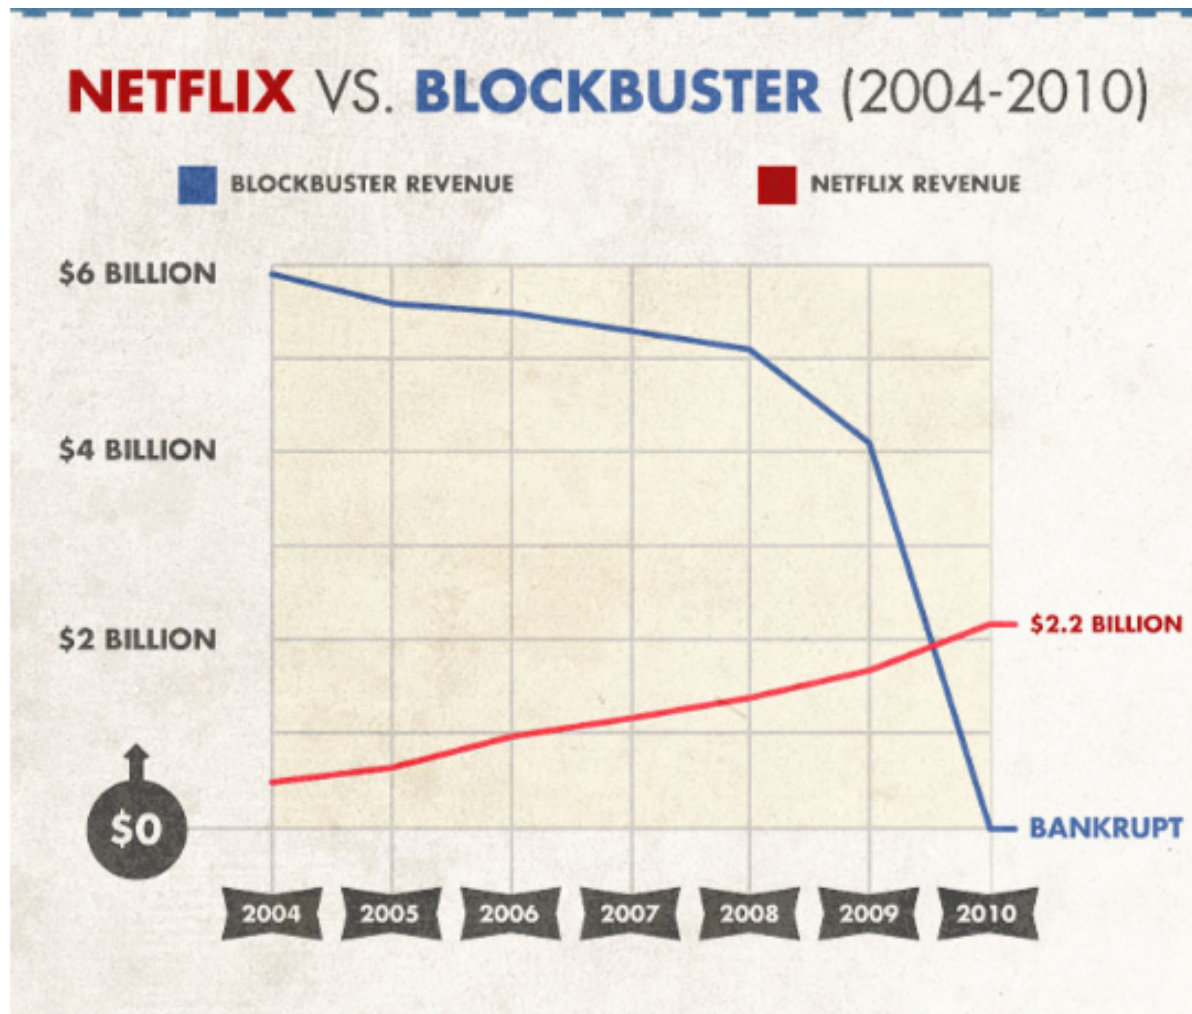
Latency in milliseconds



A study by MIT Sloan Management Review and Capgemini Consulting finds that companies now face a digital imperative: adopt new technologies effectively or face competitive obsolescence.

- October 2013

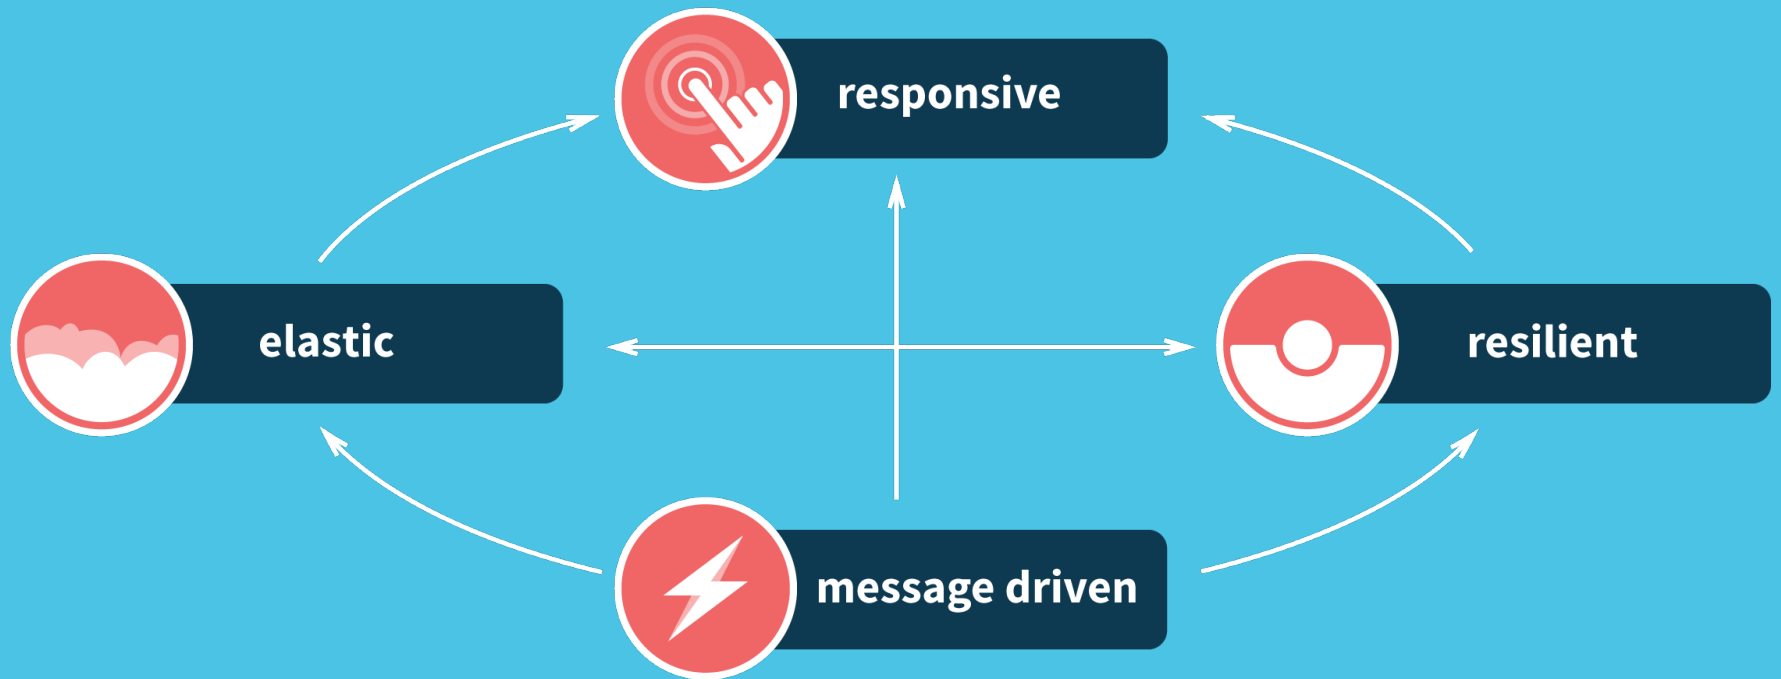
Case and Point





“In today's world, the demand for distributed systems has exploded. As customer expectations such as an immediate response, no failure, and access anywhere increase, companies have come to realize that distributed computing is the only viable solution.”
- Reactive Application Development (Manning)

Reactive Systems





“Modern applications must embrace these changes by incorporating this behavior into their DNA”.
- Reactive Application Development (Manning)

Outline

1. Introduction
2. What is Event Sourcing?
3. What are Commands?
4. What are Events?
5. What is Akka Persistence?
6. What is Akka Cluster Sharding?
7. Consistency
8. Conclusion

what is
event sourcing?



“The majority of business applications today rely on storing current state in order to process transactions. As a result in order to track history or implement audit capabilities additional coding or frameworks are required.”

- Greg Young

Current State Only??

This was **not always** the case:

- Side-effect of the adoption of **RDBMS** systems
- High performance, systems **do not** do this
- Mission critical systems **do not** do this
- RDBMS's do not do this **internally!**
- SCADA (System Control and Data Acquisition) Historian
- File control **audit** systems



“Event sourcing provides a means by which we can capture the real behavior (intent) of our users”
- Reactive Application Development (Manning)

Event Sourcing

Historical **behavior** is captured

- **Behavioral** by nature
- Convert valid **commands** into **one or more** events
- Current state **is not** persisted
- Current state is **derived**
- **Append only** store with **simple** key structure
- Designed for **distribution**



“This pattern can simplify tasks in complex domains by avoiding the requirement to synchronize the data model and the business domain”

- Reactive Application Development (Manning)

what are
commands?



command | kə`mand |

- [reporting verb] give an authoritative order: [with obj. and infinitive]

Commands

Commands are about **behavior** rather than data centrality.

Commands are a **request** of the system to perform a **task** or **action**. They follow a **VerbNoun** format, for example:

```
case class RegisterClient(id: String, . . .)
case class ChangeClientLocale(id: String, expVer: Long, . . .)
```

Commands

- Commands are **imperative**
- They are requests to **mutate state**
- An action one **would like** to take
- Transfer as **messages** not DTO's
- Implies **task-based UX**

Commands

- Conceptually, performing **task**
- **Not** data edits, rather behavioral **request**
- Can be **rejected**
- They do not **expose** internal state
- Greatly **simplified** repository layer
- Single command can = **multiple** events

what are
events?



event | i`vent |

noun

- a thing that happens, especially one of importance

Events

Events are **Indicative** in nature. They serve as a sign or **indication** that something has **happened**.

As such, they are **immutable** and cannot be **rejected**. They follow a **NounVerb** format, for example:

```
case class ClientRegistered(id: String, ver: Long, . . .)
case class ClientLocaleChanged(id: String, ver: Long, . . .)
```

Events

- **Atomic** by nature
- Record of state **change**
- **VerbNoun** implies behavior
- **Immutable**
- Natural, **verifiable** audit log
- Cannot be **rejected**

Canonical Example

One of the best ways to understand event sourcing is to look at the **canonical** example, a bank account register.

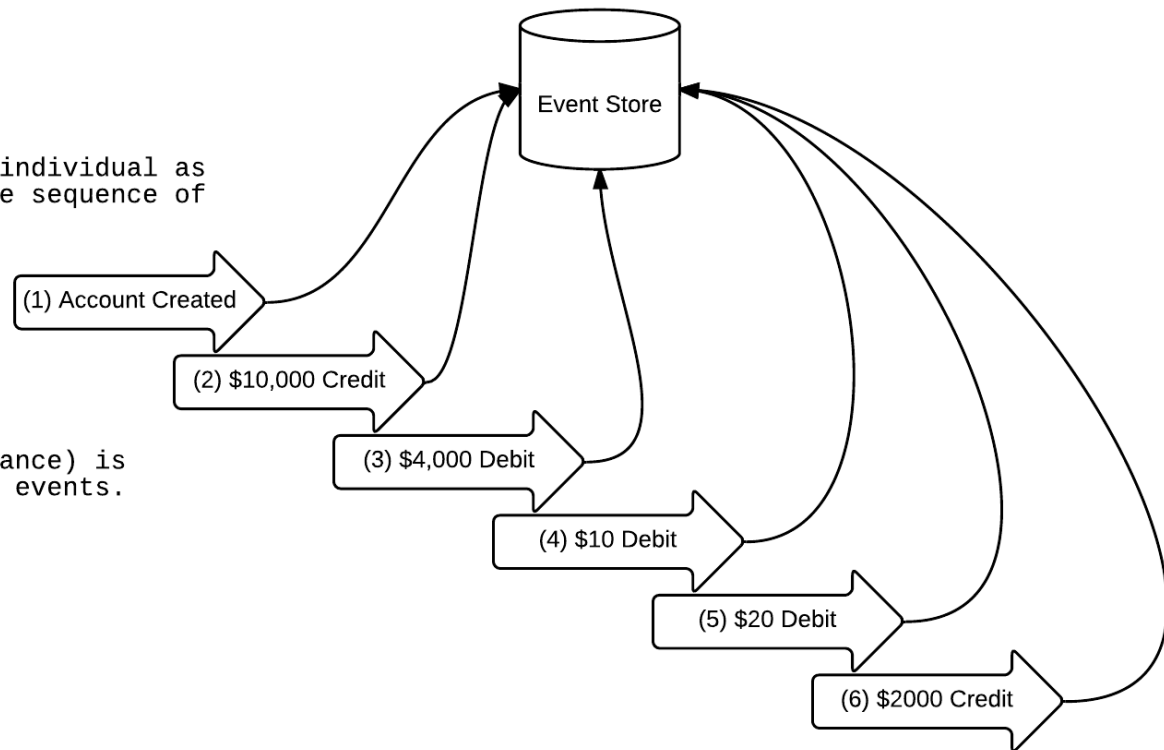
Date	Comment	Change	Balance
7/1/2014	Deposit from 3300	+ 10,000.00	10,000.00
7/3/2014	Check 001	- 4,000.00	6,000.00
7/4/2014	ATM Withdrawal	- 3.00	5,997.00
7/11/2014	Check 002	- 5.00	5,992.00
7/12/2014	Deposit from 3301	+ 2,000.00	7,992.00

Canonical Example

Each event is persisted individual as a delta. Following is the sequence of events:

1. Account Created
2. \$10,000 Credit
3. \$4,000 Debit
4. \$10 Debit
5. \$20 Debit
6. \$2,000 Credit

NOTE: Current state (balance) is derived by replaying all events.



Canonical Example

- We persist each transaction as an **independent** event
- To calculate the balance, the **delta** of the current transaction is **applied** to the last known value
- We have a **verifiable** audit log that can be **reconciled** to ensure **validity**
- The current balance at **any** point can be **derived** by **replaying** all the **transactions** up to that point
- We have captured the **real** intent of **how** the account holder manages their finances

Outline

1. Introduction
2. What is Event Sourcing?
3. What are Commands?
4. What are Events?
- 5. What is Akka Persistence?**
6. What is Akka Cluster Sharding?
7. Consistency
8. Conclusion

what is
akka persistence?

Akka Persistence

- **PersistentActor** model can **persist** internal state
- Underlying semantics use **event sourcing**
- **Append-only** store w/ **community** plugins
- **State** recovered by **replaying** events/snapshots
- **Point-to-point** communication with **at-least-once** message delivery
- **Identity** or **Identity and State!**

PersistentActor

- Persistent, **stateful** actor model
- Reacts in a **thread-safe** manner
- Supports **command** sourcing
- Supports **event** sourcing
- Implicit **replay** optimization
- **Recovers** the state from journal/snapshots

Journal

- Event sourced **append-only** storage model
- Application **controls** message journaling
- Journal implementation is **pluggable**
- The **default** journals to the filesystem
- **Replicated journals @ Community Plugins**

Snapshots

- A snapshot stores a “**moment-in-time**”
- Snapshots **internal state** of the actor
- Used for **optimizing** recovery times
- Snapshot implementation is **pluggable**.
- The **default** snapshots to the filesystem.
- **Replicated** snapshots @ Community Plugins

receiveCommand

```
class Client extends PersistentActor {  
  
  . . .  
  
  val receiveCommand: Receive = { //<- process commands  
  
    case cmd: RegisterClient => validateRegistration(cmd) fold (  
  
      f => sender ! f,  
  
      s => persist(Event) { e =>  
  
        state = state.update(e)  
  
        // side effects go here  
  
        . . .  
  
      }  
    }  
  }  
}
```

Identity and State

```
object Client {  
  . . .  
  private def empty: Client = Client()  
  private case class State(c: Client) {  
    def update(e: Event): State = e match {  
      . . .  
    }  
  }  
}  
  
class Client extends PersistentActor {  
  var state = State(empty) //<- mutable state OK!  
  . . .  
}
```

receiveRecover (resilience)

```
class Client extends PersistentActor {  
  . . .  
  val receiveRecover: Receive = {  
    case e: Event => e match {  
      case evt: ClientRegistered =>  
        state = state.update(evt)  
        // there should be no side effects here  
        . . .  
    }  
    case SnapshotOffer(_, snapshot: Client) => state = snapshot  
  }  
}
```

what is
akka cluster sharding?

Akka Cluster Sharding

- **Identity** distribution across several nodes
- One machine is **not enough**, cluster required
- Naturally **elastic & resilient**
- Actor **activation & passivation**
- Messages sent to **shard** (think proxy)
- Physical location of **actor** is managed

Shard Region

- ShardRegion actor on **each** cluster **node**
- Entry point for **entity** identification
- **Extracts** identifier from **incoming** messages
- If actor location **unknown**, asks ShardCoordinator

Shard Coordinator

- ShardCoordinator **decides** (1st message):
 - which **ShardRegion** that owns the shard
- **Subsequent** messages:
 - delivered to the **target** destination
 - **without** involving the ShardCoordinator

Shard Region

• • •

```
val clientRegion: ActorRef = ClusterSharding(system).start(  
  typeName = Client.shardName,  
  entryProps = Some(Client.props),  
  idExtractor = Fellow.idExtractor,  
  shardResolver = Client.shardResolver)
```

• • •

```
val cmd = ChangeClientName("123", "Jason", expVer=4)  
clientRegion ! cmd
```

Outline

1. Introduction
2. What is Event Sourcing?
3. What are Commands?
4. What are Events?
5. What is Akka Persistence?
6. What is Akka Cluster Sharding?
- 7. Consistency**
8. Conclusion

what is
consistency?



consistency | kən'sistənsē |
noun

- conformity in the application of something, typically necessary for the sake of logic; accuracy or fairness



“Consistency is often taken for granted when designing traditional monolithic systems as you have tightly coupled services connected to a centralized database”
- Reactive Application Development (Manning)

Strong Consistency

Traditional monolithic systems default to Strong Consistency as there is only **one path** to the data store for a given service and that path is **synchronous** in nature.

- All **accesses** are available to all **processes**
- All **accesses** are seen in the same **sequential** order

Strong Consistency

In distributed computing, however, this is **not the case**. By design, distributed systems are **asynchronous** and **loosely coupled** and rely on patterns such as atomic shared memory systems and distributed data stores achieve Availability and Partition Tolerance

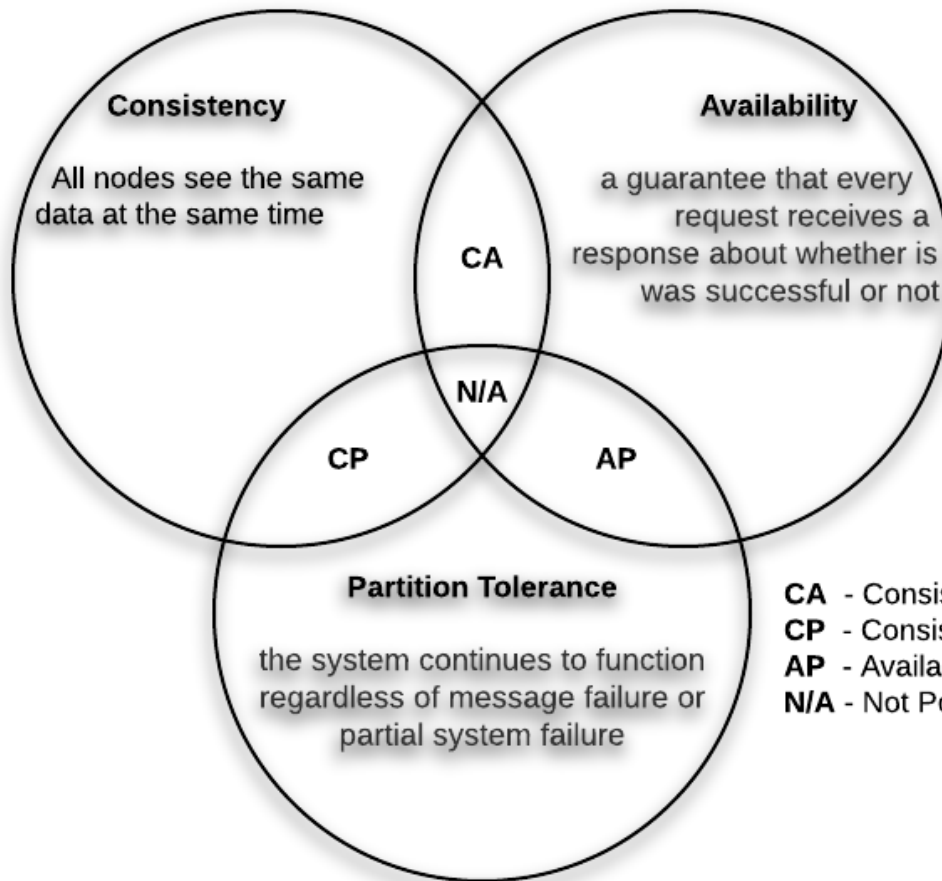
Therefore, strongly consistent systems are not distributable **as a whole contiguous system** as identified by the CAP theorem.

CAP Theorem

In Theoretical Computer Science, CAP Theorem, also known as Brewer's Theorem, states that its **impossible** in Distributed Systems to **simultaneously** provide **all three** of the following **guarantees**:

- Consistency - all **nodes** see the **same** data at the **same** time
- Availability - a guarantee that **every** request receives a **response** about whether successful or not
- Partition Tolerance - the system **continues** to **function** regardless of **message** failure or **partial system** failure

CAP Theorem



CA - Consistent and Available
CP - Consistent and Partition Tolerant
AP - Available and Partition Tolerant
N/A - Not Possible



“In distributed computing, a system supports a given consistency model if operations follow specific rules as identified by the model. The model specifies a contractual agreement between the programmer and the system, wherein the system guarantees that if the rules are followed, memory will be consistent and the results will be predictable.”

- Wikipedia

Eventual Consistency

Eventual consistency is a consistency model used in distributed computing that **informally guarantees** that, if no new updates are made to a given data item, **eventually** all accesses to that item will return the last **known** value.

- **Pillar** of distributed systems
- Often under the moniker of **optimistic replication**
- Matured in the **early** days of **mobile** computing

Eventual Consistency

A system that has **achieved** eventual consistency is often said to have **converged**, or achieved replica convergence.

- While stronger models, like **linearizability** (Strong Consistency) are **trivially** eventually consistent, the converse does not hold.
- **Eventually Consistent** services are often classified as as **Basically Available Soft state Eventual consistency** semantics as opposed to a more traditional **ACID** (Atomicity, Consistency, Isolation, Durability) guarantees.

Causal Consistency

Causal consistency is a **stronger** consistency model that **ensures** that the operations processes in the order expected.

More precisely, partial order over operations is **enforced** through **metadata**.

- If operation **A** occurs before operation **B**, then any data center that sees operation **B** must see operation **A** first.

There are three rules that define potential causality.

Causal Consistency (3 Rules)

- **Thread of Execution:** If A and B are two operations in a single thread of execution, then $A \rightarrow B$ if operation A happens before B .
- **Reads-From:** If A is a write operation and B is a read operation that returns the value written by A , then $A \rightarrow B$.
- **Transitivity:** For operations A , B , and C , if $A \rightarrow B$ and $B \rightarrow C$, then $A \rightarrow C$. Thus the causal relationship between operations is the **transitive closure** of the first two rules.

what is
conflict resolution?



resolution | rezə'ləōōSHən |

noun

- a firm decision to do or not to do something

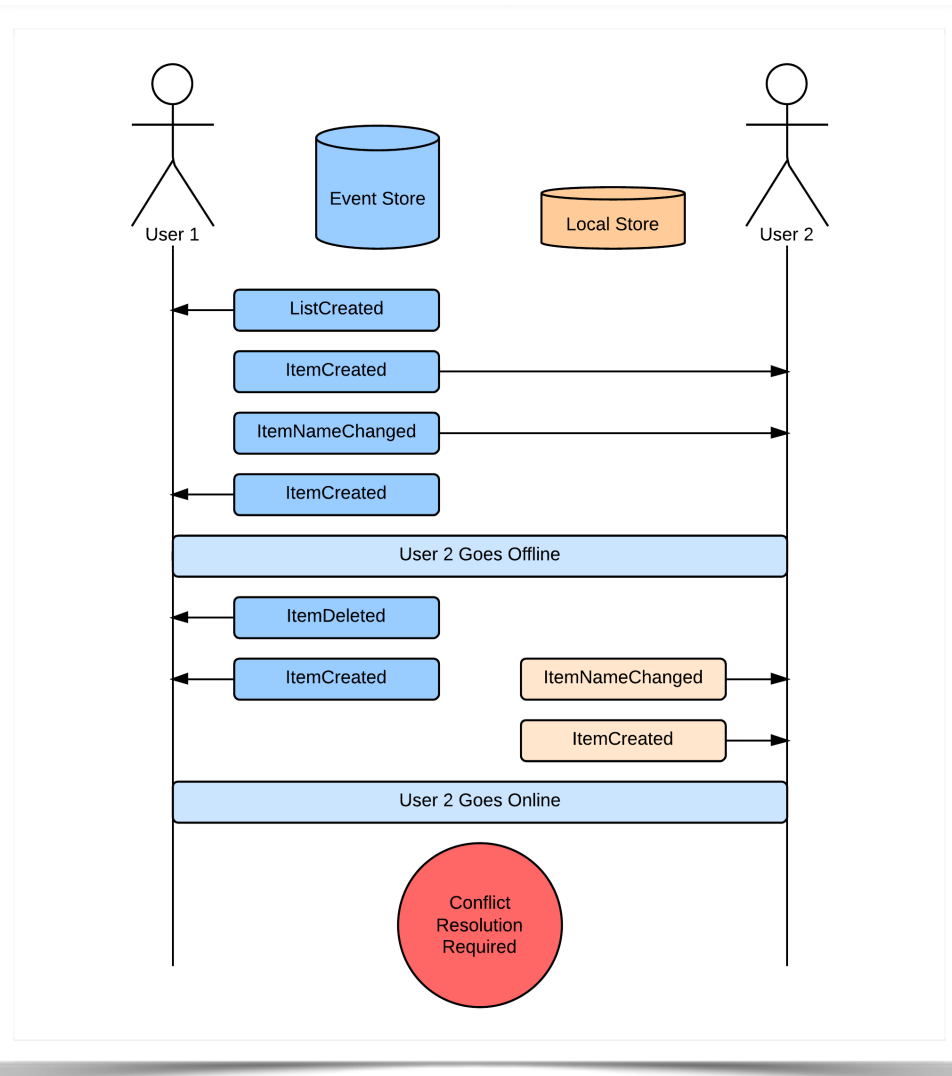


“In order to ensure the convergence of replicated data, a reconciliation between the distributed copies is required.

This process, often known as [anti-entropy], requires versioning semantics to as part of the data”

- Wikipedia

Conflict Resolution



Conflict Resolution

The recommended way to solve is the problem for the command side of CQRS, is by **embedding** into the data structure a simple **metadata** attribute, **version number**.

- Known as **current state** versioning
- Does **command** version = **current state** version?
- If they are not equal, the command is **rejected**
- **First writer wins**.


Conflict Resolution

```
object Client {  
  def requireVersion[C <: EventableCommand]  
    (c: Client, cmd: C): Either[ErrorMsg, C] =  
    if(cmd.expVer == c.ver)  
      Right(cmd)  
    else  
      Left(ErrorMsg(List("Expected version mismatch")))  
  . . .  
}
```

Outline

1. Introduction
2. What is Event Sourcing?
3. What are Commands?
4. What are Events?
5. What is Akka Persistence?
6. What is Akka Cluster Sharding?
7. Consistency
8. Conclusion

Reactive Application Development



Questions?

Distributed ES with Akka Persistence

Duncan K. DeVore

Typesafe

@ironfish

